

Sweetgreen: RPA-Candidate Scan


Automate the rule, route the exception — where a bot belongs on Sweetgreen's order-and-payment flow

IE-GY 9113B • Systems Integration: From ERP to Agentic AI

Week 4 • Board Memo §4 (Automation Journey) • Storyboard S6 • Milestone 1: Diagnose

Where each step actually belongs

Sweetgreen order & payment flow · IE-GY 9113B Lab 1 process map, sharpened Week 4



★ control point — data is born · ◆ decision point — routing choice · - - - rejection / rework path · ⚡ value leak (Four Lenses)
Matches the control register and diagnostic points in the Lab 1 process map and the Week 4 RPA-candidate scan — a mirror, not a template.

Two candidates, two tools

Repetitive + rules-based + stable, or it isn't a candidate. Both of Sweetgreen's diagnostic points from Lab 1 pass — into different tools.

D2 — Feasible? (live vs. cached availability)

→ RPA • Lab 2 (Make)

Repetitive: every order passes through it. Rules-based: approve when live and cached availability agree and the customization is supported, else hold — a complete rule once the freshness gate is folded in. Stable: the rule and its two source systems (Olo, Crunchtime) don't change often.

Thread: Operations & Supply Chain, primarily — also touches Cash (fewer wrongly-accepted orders) and Customer (fewer holds).

Prize: 6 of 30 synthetic orders (20%) hit the mismatch and enter exception recovery — \$44 modeled exposure, 18.7 min late on average.

Unattended on the routine path; the human stays at the exception queue, every time the gate disagrees.

D3 — Recovery agreed? (staff assess options)

→ AI agent • Lab 3, Week 9

Repetitive: every exception reaches it. Rules-based: no — two reasonable staff members could offer different resolutions to the same customer; depends on urgency, prep ETA, and consent. Fails the rules test outright.

This is the judgment step: an assistant can recommend a recovery option for manager review, but it can't make the customer's choice.

Not this week's build target — previewed now, built in Lab 3 once intelligent automation (Week 7) and agent design (Week 8) are covered.

What we rejected, and why

A scan that only says yes hasn't done the diagnostic work. Three steps, three different reasons to say no.

D1 — Order acceptance beyond payment

Repetitive and probably stable, but the map doesn't yet state the full acceptance rule. Automating an underspecified rule risks locking in an implicit policy — a mapping gap, not an automation call yet.

Quality gate — does the meal match the order?

Passes the rules test on paper, but fails the RPA pattern itself: the input is a physical, visually verified meal, not screen-readable system data. Kept as a routing decision, not elevated to a control point.

The stock-accuracy root cause behind the mismatch

Automating around a stale cache (checking it faster) would pave the cow path. The fix is closing the sync gap between Olo and Crunchtime — a redesign question, not an RPA question.

D2, costed and governed

~10

staff-hours a day, order of magnitude, at an assumed 300 orders/store/day and a 20% mismatch rate — not sourced, validate before board use

The number, order-of-magnitude

10 min manual reconciliation per affected order (60 min / 6 orders in the 30-order sandbox). Labor-removed side of payback; the Make build + licence side is Lab 2's number, not yet known. This is a shape, not a commitment — validate against real per-store volume and mismatch rates.

Attended or unattended?

Unattended, with a human checkpoint

Unattended on the routine path — no human needed when live and cached availability agree. The human stays exactly where the map already routes them: at the exception queue, every time the gate disagrees. Not case-by-case on the happy path — every time on the path that needs judgment.

Human oversight & workforce note

A narrower claim than “redeployment”

Every freshness disagreement routes to a person, never auto-resolved — guessing which source is right is exactly the silent-wrong-answer risk a confident bot creates.

This candidate removes a manual copy-and-compare check, not a role. Freed minutes return to staff already doing exception handling. A fuller workforce answer, as automation scales past this one candidate, is Board Memo \$7's job.

WHERE THIS LANDS

The scan is the argument. The build is Lab 2.

Board Memo §4

Automation Journey — this scan is the evidence

Slide S6

Automation-opportunity slide, still inside Milestone 1: Diagnose

Lab 2, Week 6

Build the D2 freshness-and-feasibility gate in Make, then break it

Lab 3, Week 9

Build the D3 recovery-assist agent, previewed here

Full detail: wiki/project/RPA-Candidate-Scan-Sweetgreen.md